

Fibonacci Numbers: A Deep Dive

Stanislav Ostapenko

October 14, 2025

Abstract

This article explores the Fibonacci sequence from basic implementations to advanced mathematical techniques. We begin with a naive recursive method, highlighting its exponential complexity and stack limitations in Java. Using a C++ JVMTI agent, we analyze JVM stack frames to understand `StackOverflowException`. We address Java's `long` type limitations with `BigDecimal` for large numbers. Optimization techniques like memoization and dynamic programming are introduced to improve performance. We derive Binet's formula using formal power series and explore matrix exponentiation for logarithmic-time computation. Finally, we discuss real-world applications, including the golden ratio, algorithms, and financial modeling.

Contents

1	Recursion and Mathematical Induction	5
2	Naive Recursion	6
2.1	Algorithm in Java	6
2.2	Binet's formula	6
2.2.1	Intuitive Explanation	6
2.2.2	Formal Derivation	7
2.3	Time Complexity (Big \mathcal{O})	8
2.4	Empirical Validation of Time Complexity	8
2.5	Introduction to JVM Memory Structures	12
2.6	Method Execution in Java	14
2.6.1	Core concepts	14
2.6.2	Pass-by-Value in Java	15
2.6.3	Tail Recursion	16
2.7	Depth and StackOverflow	16
2.8	Recursion Tree	17
2.9	Space Complexity	21
2.9.1	Call Stack Analysis	21
2.9.2	Clarification on Exponential Misconception	21
2.10	JVM Debugger view	21
2.11	Conclusion	30
3	Optimizing Recursion	31
3.1	Memoization	31
3.2	Dynamic Programming	31
3.3	Linear algebra and Fibonacci	32
3.3.1	Matrices and Transformations	32
3.3.2	Algorithm implementation in Java	33
3.3.3	Time and space complexity	33
3.4	Computer Algebra Systems	34
3.5	Binet formula in real life	34
3.6	Conclusion	35
4	Limitations of Primitive Types for Large Fibonacci Numbers	36
4.1	Showcase : from int to double	36
4.2	When double goes wild	38
4.2.1	$0.1 + 0.2 \neq 0.3$	38
4.2.2	IEEE 754 Representation	39
4.2.3	Calculate as machines	39
4.2.4	Conclusion	40
4.3	Handling Large Numbers with BigDecimal	40
5	Real-World Applications	41
5.1	Fibonacci heaps for Dijkstra's algorithm optimization	41
5.2	Fibonacci retracement levels in stock market analysis	41
5.3	Some of pseudorandom number generators	41

Listings

1	Naive Recursive Fibonacci in Java	6
	<code>./recursive-time-exec/RecursiveGrowthDemonstrator.java</code>	8
	<code>./recursive-time-exec/fib-exec-time-chart.py</code>	9
	<code>./recursive-time-exec/empiristic-formula-for-time-complexity.py</code>	11
2	Testing Recursion Depth in Java	16
3	Running RecursionDepth	17
4	Minimal Debugger	22
5	Debugger Target	25
6	Call frames in Java Debugger	26
7	Memoized Fibonacci	31
8	Array-Based DP Fibonacci	31
9	Space-Optimized DP Fibonacci	31
10	Matrix Exponentiation for Fibonacci	33
11	Fibonacci in Maple	34
12	Running Maple Script	34
	<code>./data-types-limitations/LargestExactFibonacci.java</code>	36
13	Fibonacci with BigDecimal	40

$$\mathcal{O}(1) = \mathcal{O}(\text{yeah})$$

$$\mathcal{O}(\log n) = \mathcal{O}(\text{nice})$$

$$\mathcal{O}(n) = \mathcal{O}(\text{k})$$

$$\mathcal{O}(n^2) = \mathcal{O}(\text{my})$$

$$\mathcal{O}(2^n) = \mathcal{O}(\text{no})$$

$$\mathcal{O}(n!) = \mathcal{O}(\text{mg})$$

$$\mathcal{O}(n^n) = \mathcal{O}(\text{sh}^*\text{t!})$$

1 Recursion and Mathematical Induction

The Fibonacci sequence, defined as $F_n = F_{n-1} + F_{n-2}$ with $F_0 = 0$ and $F_1 = 1$, is a fundamental concept in mathematics and computer science. Introduced by Leonardo of Pisa in 1202, it appears in nature (e.g., spiral patterns), algorithms (e.g., Fibonacci heaps), and number theory. We'll start our journey from naive recursion to advanced techniques, analyzing their computational complexity and practical limitations.

The concepts of recursion and mathematical induction are closely intertwined, as both rely on solving problems by breaking them down into smaller instances and establishing a base case. Below, we explore their relationship through their structural similarities and shared principles, with a particular emphasis on the role of the base case.

In mathematical induction, the base case establishes the truth of a statement for an initial value. In recursion, the base case is equally critical, as it defines the condition under which the recursive process terminates, returning a specific value without further recursive calls. The base case prevents infinite recursion and provides a foundation for building solutions to larger instances. Without a well-defined base case, a recursive function would continue indefinitely, leading to errors such as stack overflow.

For example, in a recursive factorial function, the base case is typically defined for $n = 0$ or $n = 1$, returning 1. This ensures that the recursion stops at a known value, allowing the algorithm to compute results for larger inputs by building on this foundation.

The base case is the cornerstone of both recursion and mathematical induction:

- **Termination:** In recursion, the base case ensures the process stops, preventing infinite recursion. Without it, the function would attempt to compute values for invalid inputs (e.g., negative numbers) or never terminate.
- **Correctness:** The base case aligns with the mathematical definition of the problem, ensuring accurate results. For factorial, $0! = 1$ and $1! = 1$ are standard definitions.
- **Foundation:** It provides a starting point that recursive calls or inductive steps rely on to build the solution or proof.

Both recursion and mathematical induction rely on the principle of breaking down a problem into simpler components:

- **Mathematical induction** proves a statement for all cases by starting with a base case and using the inductive step to cover all subsequent cases.
- **Recursion** computes a result by solving smaller instances of the same problem, reducing it to the base case.

Recursion and mathematical induction share a fundamental approach: solving or proving something by reducing it to simpler cases, anchored by a well-defined base case. The base case is essential for termination, correctness, and providing a foundation for building solutions or proofs. While induction is a proof technique, recursion is its practical counterpart in programming, with the base case playing a pivotal role in ensuring both processes succeed.

2 Naive Recursion

2.1 Algorithm in Java

The simplest approach to compute Fibonacci numbers is recursion, following the sequence's definition.

Time Complexity: $T(n) = \mathcal{O}(2^n)$

Space Complexity: $M(n) = \mathcal{O}(n)$ (due to call stack depth)

```
1 public static int fib0(int n) {  
2     if (n == 0) return 0;  
3     if (n == 1) return 1;  
4     return fib0(n - 1) + fib0(n - 2);  
5 }
```

Listing 1: Naive Recursive Fibonacci in Java

This method is intuitive but inefficient due to redundant calculations, forming a binary recursion tree with approximately 2^n nodes.

2.2 Binet's formula

2.2.1 Intuitive Explanation

Let us think of the Fibonacci sequence not as a list of numbers, but as the sequence of coefficients of a power series. In other words, we define a generating function

$$F(x) = F_0 + F_1x + F_2x^2 + F_3x^3 + \dots,$$

where each coefficient corresponds to a Fibonacci number.

Since the Fibonacci sequence satisfies the recurrence relation

$$F_n = F_{n-1} + F_{n-2},$$

we can express this recurrence in terms of $F(x)$ itself. To do that, consider how the series looks when multiplied by x and x^2 :

$$\begin{cases} F(x) = F_0 + F_1x + F_2x^2 + F_3x^3 + \dots, \\ xF(x) = F_0x + F_1x^2 + F_2x^3 + F_3x^4 + \dots, \\ x^2F(x) = F_0x^2 + F_1x^3 + F_2x^4 + F_3x^5 + \dots. \end{cases}$$

Now, if we take $F(x) - xF(x) - x^2F(x)$, all the shifted terms cancel out due to the recurrence relation, leaving only the initial conditions:

$$F(x) - xF(x) - x^2F(x) = F_0 + (F_1 - F_0)x.$$

Assuming $F_0 = 0$ and $F_1 = 1$, we obtain

$$F(x) = \frac{x}{1 - x - x^2}.$$

The denominator here encodes the same recurrence that defines Fibonacci numbers. To understand the structure of $F(x)$, we factor the quadratic polynomial:

$$1 - x - x^2 = (1 - \varphi x)(1 - \psi x),$$

where

$$\varphi = \frac{1 + \sqrt{5}}{2}, \quad \psi = \frac{1 - \sqrt{5}}{2}.$$

By the method of partial fractions, we can decompose $F(x)$ as

$$F(x) = \frac{A}{1 - \varphi x} + \frac{B}{1 - \psi x}.$$

Each term now has a familiar geometric series form:

$$\frac{1}{1 - rx} = 1 + rx + r^2x^2 + r^3x^3 + \dots,$$

so we can directly read off the coefficients as powers of φ and ψ .

This is why we deliberately rewrite $F(x)$ in such a form: it allows us to transform an abstract recurrence relation into a closed analytic expression. Eventually, by equating coefficients of x^n , we recover the celebrated Binet formula:

$$F_n = \frac{\varphi^n - \psi^n}{\sqrt{5}}.$$

2.2.2 Formal Derivation

We start from the Fibonacci recurrence

$$F_0 = 0, \quad F_1 = 1, \quad F_n = F_{n-1} + F_{n-2} \quad (n \geq 2).$$

Defining the generating function

$$F(x) = \sum_{n=0}^{\infty} F_n x^n.$$

Then

$$xF(x) = \sum_{n=0}^{\infty} F_n x^{n+1}, \quad x^2F(x) = \sum_{n=0}^{\infty} F_n x^{n+2}.$$

Applying the recurrence relation

$$F(x) - xF(x) - x^2F(x) = F_0 + (F_1 - F_0)x + \sum_{n=2}^{\infty} (F_n - F_{n-1} - F_{n-2})x^n.$$

Since $F_n - F_{n-1} - F_{n-2} = 0$ for all $n \geq 2$, we get

$$F(x) - xF(x) - x^2F(x) = x.$$

Hence,

$$F(x) = \frac{x}{1 - x - x^2}.$$

Factorization and substitution. Let

$$1 - x - x^2 = (1 - \varphi x)(1 - \psi x),$$

where

$$\varphi = \frac{1 + \sqrt{5}}{2}, \quad \psi = \frac{1 - \sqrt{5}}{2}.$$

Then

$$F(x) = \frac{x}{(1 - \varphi x)(1 - \psi x)} = A \frac{x}{1 - \varphi x} + B \frac{x}{1 - \psi x}.$$

Solving for constants. Multiplying both sides by $(1 - \varphi x)(1 - \psi x)$:

$$x = A(1 - \psi x) + B(1 - \varphi x) = (A + B) - (\psi A + \varphi B)x.$$

Matching coefficients gives

$$A + B = 0, \quad \varphi B + \psi A = -1.$$

Solving:

$$A = \frac{1}{\varphi - \psi}, \quad B = -\frac{1}{\varphi - \psi}.$$

Geometric series expansion

$$\frac{1}{1 - rx} = \sum_{n=0}^{\infty} r^n x^n.$$

Thus,

$$F(x) = \frac{1}{\varphi - \psi} \left(\frac{x}{1 - \varphi x} - \frac{x}{1 - \psi x} \right) = \frac{1}{\varphi - \psi} \sum_{n=1}^{\infty} (\varphi^n - \psi^n) x^n.$$

Extracting coefficients. The coefficient of x^n gives

$$F_n = \frac{\varphi^n - \psi^n}{\varphi - \psi} = \frac{1}{\sqrt{5}} \left[\left(\frac{1 + \sqrt{5}}{2} \right)^n - \left(\frac{1 - \sqrt{5}}{2} \right)^n \right].$$

$$F_n = \frac{\varphi^n - \psi^n}{\sqrt{5}}$$

This is the closed form known as **Binet's formula**.

2.3 Time Complexity (Big \mathcal{O})

The recursive algorithm generates a binary recursion tree, where each node for $n \geq 2$ spawns two child nodes: $F(n-1)$ and $F(n-2)$. The total number of function calls corresponds to the number of nodes in the recursion tree. For a given n , the tree has a depth of approximately n , and the number of nodes grows exponentially. The recurrence relation for the number of operations $T(n)$ is:

$$T(n) = T(n-1) + T(n-2) + O(1),$$

where $O(1)$ accounts for the constant-time addition operation. The base cases are:

$$T(0) = O(1), \quad T(1) = O(1).$$

This recurrence is similar to the Fibonacci sequence itself. The number of nodes is approximately $2F(n) - 1$, where $F(n) \approx \phi^n / \sqrt{5}$, and $\phi = \frac{1+\sqrt{5}}{2} \approx 1.618$ is the golden ratio. Thus, the time complexity is:

$$T(n) = O(\phi^n) \approx O(1.618^n).$$

2.4 Empirical Validation of Time Complexity

Let's calculate execution time of first 50 Fibonacci numbers. Also, save exec results in CSV file further for analysis.

```
1 import java.io.FileWriter;
2 import java.io.IOException;
3 import java.io.PrintWriter;
4 import java.util.concurrent.TimeUnit;
5
6 public class RecursiveGrowthDemonstrator {
7     public static long fibonacciRecursive(int n) {
8         if (n <= 1) {
9             return n;
10        }
11        return fibonacciRecursive(n - 1) + fibonacciRecursive(n - 2);
12    }
13 }
```



```

14 public static void main(String[] args) {
15     String filename = "fibonacci_data.csv";
16     int last_n = 50;
17     try (PrintWriter writer = new PrintWriter(new FileWriter(filename))) {
18         // Write the CSV file header
19         writer.println("n,Fn,time_sec");
20         for (int n = 1; n <= last_n; n++) {
21             long startTime = System.nanoTime();
22             long result = fibonacciRecursive(n);
23             long endTime = System.nanoTime();
24
25             double durationSec = (double) (endTime - startTime) / 1_000_000_000.0;
26             writer.printf("%d,%d,%.10f\n", n, result, durationSec);
27             System.out.printf("F(%d) calculated in %.4f sec.\n", n, durationSec);
28         }
29     } catch (IOException e) {
30         System.err.println("Error while writing to file: " + e.getMessage());
31     }
32 }
33 }

```

To show how the time grows, let's build a chart -




Figure 1: Exponent execution time

To produce this image we use Python with some Pandas :

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import os
5
6 # --- 1. Define the filename and check for its existence ---
7 FILENAME = 'fibonacci_data.csv'
8

```

```

9  if not os.path.exists(FILENAME):
10     print(f"Error: File '{FILENAME}' not found.")
11     print("Please run the Java code first to generate the data file.")
12     exit()
13
14  # --- 2. Read and Prepare Data ---
15  try:
16     # Read the data from the CSV file
17     df = pd.read_csv(FILENAME)
18  except Exception as e:
19     print(f"Error reading CSV file: {e}")
20     exit()
21
22  # Filter out very small execution times (mostly for n < 30)
23  # as they introduce noise, focusing the graph on the exponential growth phase.
24  # We'll keep only data points where time is greater than 1 millisecond (0.001 sec).
25  df_filtered = df[df['time_sec'] > 0.001].copy()
26
27  if df_filtered.empty:
28     print("Not enough data points with significant execution time (above 0.001 sec)
29         to plot exponential growth.")
30     print("Try increasing the 'last_n' value in your Java code (e.g., to 45).")
31     exit()
32
33  # --- 3. Plotting the Exponential Growth ---
34  plt.figure(figsize=(12, 7))
35
36  # Plot the actual recursive time
37  plt.plot(df_filtered['n'], df_filtered['time_sec'],
38          marker='o', linestyle='-', color='red', label='Actual Recursive Time (sec)',
39          linewidth=2)
40
41  plt.title('Exponential Time Growth: Naive Recursive Fibonacci (O($\phi^n$))',
42          fontsize=16)
43  plt.xlabel('Fibonacci Number Index (n)', fontsize=14)
44  plt.ylabel('Execution Time (seconds)', fontsize=14)
45  plt.legend(fontsize=12)
46  plt.grid(True, which='both', linestyle='--', linewidth=0.5)
47
48  # Highlight the steep rise for visual emphasis
49  if len(df_filtered) > 5:
50     steep_start_n = df_filtered[df_filtered['time_sec'] >
51         df_filtered['time_sec'].max() * 0.1]['n'].min()
52     plt.axvline(x=steep_start_n, color='gray', linestyle=':', linewidth=1.5,
53         label=f'Steep Growth Point (n $\approx$ {steep_start_n})')
54     plt.legend(fontsize=12)
55
56  # --- 4. Adding Annotations for Educational Value ---
57
58  # Find the last calculated point
59  last_n_point = df_filtered.iloc[-1]
60  plt.annotate(
61     f'Time: {last_n_point["time_sec"]:.2f} sec\n(n={last_n_point["n"]})',
62     xy=(last_n_point['n'], last_n_point['time_sec']),
63     xytext=(last_n_point['n'] - 10, last_n_point['time_sec'] * 0.8),
64     arrowprops=dict(facecolor='blue', shrink=0.05, width=1, headwidth=8),
65     fontsize=11,

```

```

64     color='blue'
65 )
66
67 # Text explanation on the graph
68 plt.text(df_filtered['n'].min() + 2, df_filtered['time_sec'].max() * 0.9,
69         'O($\phi^n$): Time nearly multiplies by 1.618 with each increment of n.',
70         fontsize=12, bbox=dict(facecolor='yellow', alpha=0.5))
71
72 plt.savefig("exponential-time-growth.png")
73 plt.show()

```

Now we are interested in exact formula of this type of growth. To achieve our goal we'll going to use SciPy.

```

1  import pandas as pd
2  import numpy as np
3  import matplotlib.pyplot as plt
4  from scipy.optimize import curve_fit
5
6  # Define the exponential model function
7  def exponential_model(n, a, b):
8      return a * b**n
9
10 # Read data from CSV file
11 data = pd.read_csv('fibonacci_data.csv')
12
13 # Extract n and time_sec columns
14 n = data['n'].values
15 time_sec = data['time_sec'].values
16
17 # Fit the exponential model
18 popt, pcov = curve_fit(exponential_model, n, time_sec, p0=[1e-6, 1.618]) # Initial
    guess: a=1e-6, b=1.618
19 a, b = popt
20 print(f"Fitted model: T(n) = {a:.10f} * {b:.6f}^n")
21
22 # Compute predicted time values
23 predicted_time = exponential_model(n, a, b)
24
25 # Plot the results
26 plt.figure(figsize=(10, 6))
27 plt.scatter(n, time_sec, color='blue', label='Experimental data')
28 plt.plot(n, predicted_time, color='red', label=f'Model: T(n) = {a:.2e} * {b:.6f}^n')
29 plt.xlabel('n')
30 plt.ylabel('Execution time (sec)')
31 plt.title('Execution time of recursive Fibonacci algorithm')
32 plt.yscale('log') # Log scale for better visualization of exponential growth
33 plt.legend()
34 plt.grid(True)
35 plt.savefig("experimental-data-formula.png")
36 plt.show()
37
38 # Compare b to the golden ratio
39 phi = (1 + np.sqrt(5)) / 2
40 print(f"Golden ratio  $\phi$ (): {phi:.6f}")
41 print(f"Deviation of b from  $\phi$ : {abs(b - phi):.6f}")

```

Which shows us the following :

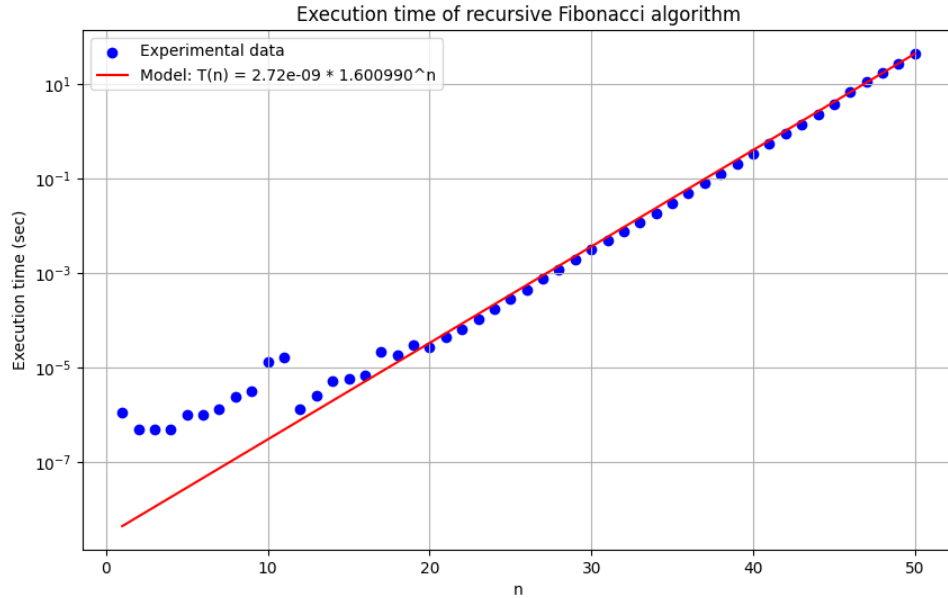


Figure 2: Exponent execution time

Axes made logarithmic for better looking chart. So, we have result of our numerical modeling

```

1 python empiristic-formula-for-time-complexity.py
2 Fitted model:  $T(n) = 0.0000000027 * 1.600990^n$ 
3 Golden ratio  $\phi()$ : 1.618034
4 Deviation of b from  $\phi$ : 0.017044

```

So, we can have exec time needed to calculate n-th fib. It grows fast and could be very big.

Table 1: Estimated execution time of recursive Fibonacci algorithm using the model $T(n) = 2.7 \times 10^{-9} \cdot 1.600990^n$

n	$T(n)$	Unit
20	3.30e-05	sec
30	3.66e-03	sec
40	4.04e-01	sec
50	4.48e+01	sec
60	1.3752	hr
70	6.3395	days
80	1.9215	yr
90	212.5851	yr
100	23519.0125	yr

So, we have kind of bad news here. To calculate $F(100)$ we'll need ≈ 25000 years. Too long to wait, soon we improve algorithm.

2.5 Introduction to JVM Memory Structures

The Java Virtual Machine (JVM) is a runtime environment that executes Java bytecode, enabling platform-independent execution of Java programs. The JVM manages memory through a structured architecture that supports dynamic allocation, thread execution, and garbage collection.

The JVM divides its memory into several regions, each serving a specific purpose in program execution. These regions are broadly categorized into **per-thread** and **shared** areas:

- **Per-Thread Areas:** Allocated for each thread to ensure isolation and manage method execution.
- **Shared Areas:** Accessible by all threads for storing objects and class metadata.

Memory management is critical for performance, as it affects allocation speed, garbage collection, and thread synchronization. The JVM's memory model is defined by the Java Virtual Machine Specification (JVMS).

The JVM's memory is organized into the following key areas:

1. **Heap:**

- A shared memory region where all objects and arrays are allocated using the **new** keyword.
- Divided into:
 - **Young Generation** (Eden and Survivor spaces): For newly created objects, managed by frequent minor garbage collections.
 - **Old Generation:** For long-lived objects, managed by less frequent major garbage collections.
 - **Metaspace** (Java 8+): Stores class metadata, replacing the Permanent Generation (pre-Java 8).
- Configurable via flags like **-Xmx** (maximum heap size) and **-Xms** (initial heap size).

2. **Java Stack:**

- A per-thread memory area that stores **call frames** for method invocations.
- Each frame contains a local variable array, operand stack, and frame data (e.g., program counter, return address).
- Size is configurable via **-Xss**. Excessive recursion can cause a **StackOverflowError**.

3. **Program Counter (PC) Register:**

- A per-thread register that holds the address of the current bytecode instruction being executed.
- Points to the current instruction in the active call frame's bytecode.

4. **Method Area:**

- A shared area that stores class metadata, including bytecode, constant pools, and method tables.
- In Java 8+, the Method Area is implemented as Metaspace, which uses native memory rather than the heap.

5. **Native Method Stack:**

- A per-thread stack for executing native methods (e.g., C/C++ code called via JNI).
- Similar to the Java stack but tailored for non-Java code.

Consider this Java code:

```
1 public class Example {
2     public static void main(String[] args) {
3         String str = new String("Hello");
4         int result = add(3, 4);
5         System.out.println(str + result);
6     }
7     public static int add(int a, int b) {
8         return a + b;
9     }
10 }
```

Memory Region	Type	Purpose
Heap	Shared	Stores objects, arrays, and class metadata (Metaspace in Java 8+)
Java Stack	Per-thread	Stores call frames for method execution
PC Register	Per-thread	Tracks current bytecode instruction
Method Area	Shared	Stores class metadata and constant pools
Native Method Stack	Per-thread	Manages native method execution

Table 2: JVM Memory Regions

1. **Heap:** The `String` object `"Hello"` is allocated in the heap's Eden space.
2. **Java Stack:** The `main` method's call frame stores the `str` reference and `args`. A new frame for `add` stores parameters `a` and `b`.
3. **PC Register:** Tracks the current bytecode instruction in `main` or `add`.
4. **Method Area:** Stores the bytecode and constant pool for `Example` class, including the `"Hello"` string literal.
5. **Garbage Collection:** After `main` ends, the `String` object may be reclaimed if no references remain.

The memory structures work together to support JVM execution:

- **Thread Isolation:** Per-thread areas (Java Stack, PC Register, Native Method Stack) ensure threads execute independently without interference.
- **Shared Resources:** The heap and Method Area (or Metaspace) allow threads to share objects and class data, requiring synchronization (e.g., `synchronized` blocks) to avoid race conditions.
- **Garbage Collection:** The garbage collector scans the heap, using references from stacks, Method Area, and static fields as roots to identify reachable objects.

Key Concepts :

- **Heap vs. Stack:** Heap stores dynamic, shared objects; stacks store method-scoped, thread-specific data.
- **Garbage Collection:** Automatically reclaims heap memory but requires careful reference management to avoid leaks.
- **Performance:** Memory size tuning (e.g., `-Xmx`, `-Xss`) impacts performance. Large heaps or stacks may slow execution or garbage collection.
- **Debugging:** Tools like `jstack` (for stacks), `jmap` (for heap), and VisualVM help analyze memory usage and diagnose issues.

2.6 Method Execution in Java

2.6.1 Core concepts

Java's method invocation and parameter passing mechanisms are central to understanding its runtime behavior in the Java Virtual Machine (JVM). Java exclusively uses **pass-by-value** for all parameter passing, impacting both iterative and recursive methods. Recursion, including **tail recursion**, interacts with the

JVM's stack and heap, while Java's lack of **tail call optimization (TCO)** affects performance in deep recursion.

- **Pass-by-Value:** The method receives a copy of the argument's value (primitive or object reference). Changes to the parameter do not affect the caller's variable.
- **Pass-by-Reference:** The method receives a reference to the original argument's memory location, so changes directly modify the caller's variable.

Java uses **pass-by-value** exclusively. For **primitive types** (e.g., `int`, `double`), the value is copied. For **object references**, the reference (not the object) is copied, allowing modification of the object's state in the heap but not reassignment of the caller's reference.

2.6.2 Pass-by-Value in Java

When a method is called, the JVM creates a call frame on the thread's Java stack, copying arguments into the frame's local variable array:

- **Primitive Types:** The value (e.g., 5 for an `int`) is copied. Modifying the parameter changes only the local copy.
- **Object References:** The reference (memory address to a heap object) is copied. Modifying the object's state (e.g., fields) affects the heap, visible to all references. Reassigning the reference (e.g., `obj = new Object()`) is local.

This behavior applies to both iterative and recursive methods, but recursion increases stack depth, risking `StackOverflowError` for deep calls.

Consider a Java program demonstrating pass-by-value in both iterative and recursive contexts:

```
1 public class Example {
2     public static void main(String[] args) {
3         int num = 5;
4         StringBuilder sb = new StringBuilder("Factorial: ");
5         modifyPrimitive(num);
6         modifyObject(sb);
7         int result = factorialTail(num, 1, sb);
8         System.out.println("num: " + num); // Outputs: num: 5
9         System.out.println("sb: " + sb + " " + result); // Outputs: sb: Factorial:
              5*4*3*2*1 120
10    }
11
12    public static void modifyPrimitive(int x) {
13        x = 10; // Modifies local copy
14    }
15
16    public static void modifyObject(StringBuilder builder) {
17        builder.append("World"); // Modifies heap object
18        builder = new StringBuilder("New"); // Local reassignment
19    }
20
21    public static int factorialTail(int n, int acc, StringBuilder log) {
22        if (n <= 1) {
23            log.append("1");
24            return acc;
25        }
26        log.append(n + "*");
27        return factorialTail(n - 1, n * acc, log); // Tail-recursive call
28    }
29 }
```

1. **Primitive (`num`)**: In `modifyPrimitive`, `x` is a copy of `num` (5). Setting `x = 10` affects only the local copy, so `num` remains 5.
2. **Object Reference (`sb`)**: In `modifyObject`, `builder` is a copy of the reference to `StringBuilder`. `builder.append("World")` modifies the heap object, affecting `sb`. Reassigning `builder = new StringBuilder("New")` is local, so `sb` retains its reference.
3. **Tail Recursion (`factorialTail`)**: Each recursive call copies `n`, `acc`, and `log`. Modifications to `log` (e.g., `log.append(n + "*")`) persist in the heap. The recursive call is the last operation, but Java creates a new frame each time, risking stack overflow for large `n`.

2.6.3 Tail Recursion

A method is **tail-recursive** if the recursive call is the final operation, with no pending computations. In languages with **tail call optimization (TCO)**, the runtime reuses the current frame, avoiding stack growth. Java's JVM does not support TCO, so each recursive call creates a new frame, copying arguments via pass-by-value. It's because JVM prioritizes general-purpose execution and accurate stack traces for debugging over TCO.

TCO support varies across languages, impacting recursion efficiency:

- **Java**: No TCO; each call adds a frame, risking `StackOverflowError`.
- **Scala**: TCO for self-recursive calls with `@tailrec`, compiling to loops on the JVM.
- **Python**: No TCO; uses iteration or trampolining for deep recursion.
- **JavaScript**: Partial TCO (e.g., Safari supports it, V8 does not).
- **Haskell**: Full TCO with lazy evaluation, ideal for recursion-heavy code.

Language	TCO Support	Workaround
Java	None	Iteration
Scala	Yes (<code>@tailrec</code>)	None needed
Python	None	Iteration
JavaScript	Partial	Iteration
Haskell	Full	None needed

Table 3: Tail Call Optimization Across Languages

2.7 Depth and `StackOverflow`

Recursive calls create stack frames in the JVM, which can lead to a `StackOverflowError`. We demonstrate this with a simple recursive program :

```

1 public class RecursionDepth {
2     private static int depth = 0;
3
4     public static void recurse() {
5         depth++;
6         recurse();
7     }
8
9     public static void main(String[] args) {
10        try {
11            recurse();
12        } catch (StackOverflowError e) {
13            System.out.println("Max recursion depth: " + depth);

```



```

14     }
15     }
16 }

```

Listing 2: Testing Recursion Depth in Java

Run with:

```

1 java -Xss1m RecursionDepth

```

Listing 3: Running RecursionDepth

It means that even without calculating something, we limited by the value of stack. Good news is that it could be increased, but we have no clue how much we need.

2.8 Recursion Tree

```

1 public class Simple {
2     static int depth = 0;
3
4     public static int fib0(int n) throws Exception{
5         depth++;
6         System.out.println("fib0(" + n + ") depth=" + depth + " frames=" +
7             Thread.currentThread().getStackTrace().length);
8
9         if (n == 0) {
10             depth--;
11             return 0;
12         }
13         if (n == 1) {
14             depth--;
15             return 1;
16         }
17
18         int result = fib0(n - 1) + fib0(n - 2);
19         depth--;
20         return result;
21     }
22
23     public static void main(String[] args) throws Exception {
24         System.out.println("Result: " + fib0(5));
25     }
26 }

```

We have to have compare depth calculation and number of frames, make conclusions.

Output :

```

1 C:\temp\fibonacci-article>java Simple
2 fib0(5) depth=1 frames=3
3 fib0(4) depth=2 frames=4
4 fib0(3) depth=3 frames=5
5 fib0(2) depth=4 frames=6
6 fib0(1) depth=5 frames=7
7 fib0(0) depth=5 frames=7
8 fib0(1) depth=4 frames=6
9 fib0(2) depth=3 frames=5
10 fib0(1) depth=4 frames=6
11 fib0(0) depth=4 frames=6
12 fib0(3) depth=2 frames=4
13 fib0(2) depth=3 frames=5

```

```
14 fib0(1) depth=4 frames=6
15 fib0(0) depth=4 frames=6
16 fib0(1) depth=3 frames=5
17 Result: 5
```

Let's make an improvement. We'll all indent to previous code according to depth

```
1 public class SimpleIdent {
2     static int depth = 0;
3
4     public static int fib0(int n) throws Exception {
5         // print with indentation
6         String indent = " ".repeat(depth);
7         System.out.println(indent + "fib0(" + n + ") depth=" + depth);
8
9         depth++;
10        int result;
11        if (n == 0) result = 0;
12        else if (n == 1) result = 1;
13        else result = fib0(n - 1) + fib0(n - 2);
14        depth--;
15
16        System.out.println(indent + "> fib0(" + n + ") = " + result);
17        return result;
18    }
19
20    public static void main(String[] args) throws Exception {
21        System.out.println("Result: " + fib0(5));
22    }
23 }
```

Output is better. BTW code is good for debugging any recursion.

```
1 fib0(5) depth=0
2   fib0(4) depth=1
3     fib0(3) depth=2
4       fib0(2) depth=3
5         fib0(1) depth=4
6           => fib0(1) = 1
7           fib0(0) depth=4
8             => fib0(0) = 0
9             => fib0(2) = 1
10          fib0(1) depth=3
11            => fib0(1) = 1
12            => fib0(3) = 2
13          fib0(2) depth=2
14            fib0(1) depth=3
15              => fib0(1) = 1
16              fib0(0) depth=3
17                => fib0(0) = 0
18              => fib0(2) = 1
19            => fib0(4) = 3
20          fib0(3) depth=1
21            fib0(2) depth=2
22              fib0(1) depth=3
23                => fib0(1) = 1
24                fib0(0) depth=3
25                  => fib0(0) = 0
26                => fib0(2) = 1
27              fib0(1) depth=2
28                => fib0(1) = 1
29            => fib0(3) = 2
30 => fib0(5) = 5
31 Result: 5
```

We can do better. Let's build a tree using dot syntax (blah-blah-blah)

```
1 import java.io.FileWriter;
2 import java.io.IOException;
3
4 public class Simple3 {
5
6   static class NodeId {
7     int id;
8     NodeId(int id) { this.id = id; }
9   }
10
11   static int idCounter = 0;
12
13   public static int fib0(int n, FileWriter fw, NodeId parent) throws IOException {
14     int myId = idCounter++;
15     fw.write(String.format("  node%d [label=\"fib(%d)\"];\\n", myId, n));
16
17     if (parent != null) {
18       fw.write(String.format("  node%d -> node%d;\\n", parent.id, myId));
19     }
20
21     int result;
22     if (n == 0) result = 0;
23     else if (n == 1) result = 1;
24     else {
```

```

25     int left = fib0(n - 1, fw, new NodeId(myId));
26     int right = fib0(n - 2, fw, new NodeId(myId));
27     result = left + right;
28 }
29
30 return result;
31 }
32
33 public static void main(String[] args) throws IOException {
34     FileWriter fw = new FileWriter("fib_tree.dot");
35     fw.write("digraph G {\n");
36     fib0(5, fw, null);
37     fw.write("}\n");
38     fw.close();
39     System.out.println("DOT file generated: fib_tree.dot");
40 }
41 }

```

Convert it to PNG :

```
1 C:\temp\>dot -Tpng fib_tree.dot -o fib_tree.png
```

And here it is :




Figure 3: Recursion tree

As you can see in Figure 3, recursion tree is displayed nicely.
Or even better

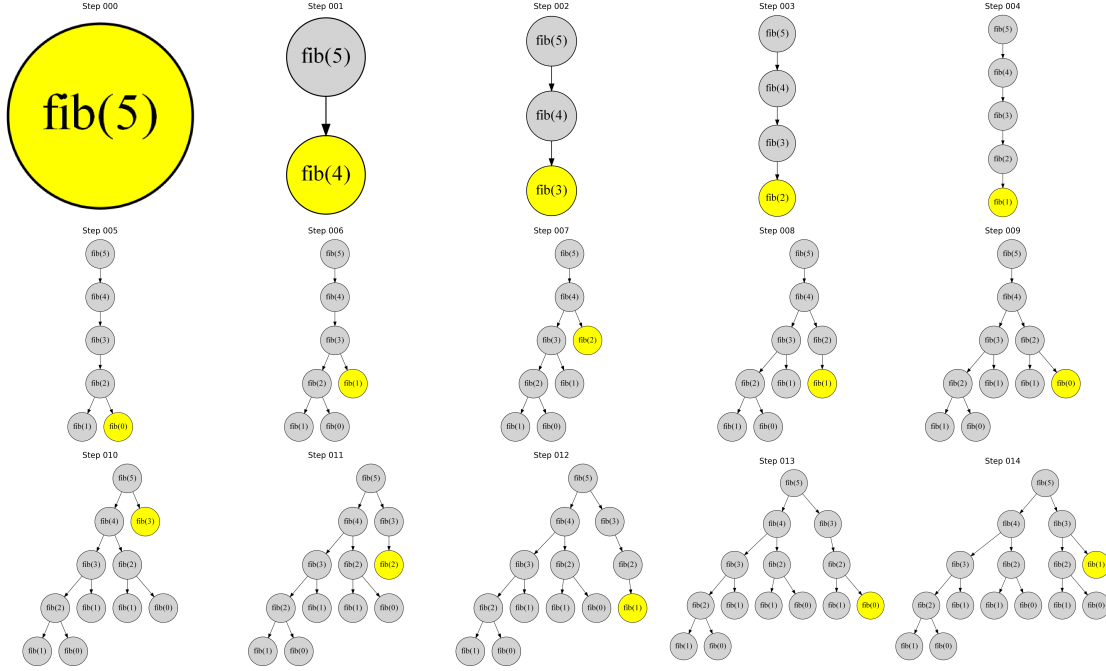


Figure 4: Recursion progress

2.9 Space Complexity

The space complexity is determined by the memory used on the call stack due to recursion. Although the recursion tree contains an exponential number of nodes ($O(\phi^n)$), not all nodes are active simultaneously. The call stack only holds the frames for the active recursive calls along a single path from the root to a leaf.

2.9.1 Call Stack Analysis

Consider the recursion tree for computing $F(n)$. The deepest path in the tree occurs when the recursion follows $F(n) \rightarrow F(n-1) \rightarrow F(n-2) \rightarrow \dots \rightarrow F(0)$, which has a depth of n . At any point, the call stack contains at most n frames, each storing a constant amount of data (e.g., the parameter n and return address).

For example: - When computing $F(n-1)$, the call for $F(n-2)$ is not yet active. - Once $F(n-1)$ is resolved, its stack frame is popped, and $F(n-2)$ is pushed onto the stack.

Thus, the maximum stack depth is n , leading to a space complexity of:

$$O(n).$$

2.9.2 Clarification on Exponential Misconception

The total number of recursive calls is exponential, which might suggest exponential memory usage. However, since only one path of the recursion tree is active at a time, the call stack grows linearly with n , not exponentially.

2.10 JVM Debugger view

```

1 import com.sun.jdi.*;
2 import com.sun.jdi.connect.*;
3 import com.sun.jdi.event.*;
4 import com.sun.jdi.request.*;
5 import java.io.IOException;
6 import java.util.*;
7
8 public class MinimalDebugger {
9     private VirtualMachine vm;
10    private EventRequestManager eventManager;
11
12    public static void main(String[] args) {
13        new MinimalDebugger().debug();
14    }
15
16    public void debug() {
17        try {
18            // 1. Підключаємося до цільової програми
19            connect();
20
21            // 2. Встановлюємо breakpoint
22            setBreakpoint();
23
24            // 3. Запускаємо програму
25            vm.resume();
26
27            // 4. Обробляємо події
28            handleEvents();
29
30        } catch (Exception e) {
31            System.err.println("Error: " + e.getMessage());
32        }
33    }
34
35    private void connect() throws IOException, IllegalConnectorArgumentsException {
36        System.out.println("Connecting to target JVM...");
37
38        VirtualMachineManager vmm = Bootstrap.virtualMachineManager();
39        AttachingConnector connector = vmm.attachingConnectors().stream()
40            .filter(c -> c.transport().name().equals("dt_socket"))
41            .findFirst()
42            .orElseThrow(() -> new RuntimeException("Socket connector not found"));
43
44        Map<String, Connector.Argument> args = connector.defaultArguments();
45        args.get("hostname").setValue("localhost");
46        args.get("port").setValue("5005");
47
48        vm = connector.attach(args);
49        eventManager = vm.eventRequestManager();
50
51        System.out.println("Connected to: " + vm.name());
52    }
53
54    private void setBreakpoint() {
55        // Чекаємо завантаження класу
56        ClassPrepareRequest classPrepareRequest =
57            eventManager.createClassPrepareRequest();
58        classPrepareRequest.addClassFilter("FibonacciTarget");
59    }

```

```

58     classPrepareRequest.enable();
59 }
60
61 private void handleEvents() throws InterruptedException {
62     EventQueue queue = vm.eventQueue();
63
64     while (true) {
65         EventSet eventSet = queue.remove();
66
67         for (Event event : eventSet) {
68             if (event instanceof ClassPrepareEvent) {
69                 handleClassPrepare((ClassPrepareEvent) event);
70             } else if (event instanceof BreakpointEvent) {
71                 handleBreakpoint((BreakpointEvent) event);
72             } else if (event instanceof VMDeathEvent) {
73                 System.out.println("Target VM terminated");
74                 return;
75             }
76         }
77
78         eventSet.resume();
79     }
80 }
81
82 private void handleClassPrepare(ClassPrepareEvent event) {
83     ReferenceType clazz = event.referenceType();
84     System.out.println("Class loaded: " + clazz.name());
85
86     // Встановлюємо breakpoint напочатку main методу
87     try {
88         Method mainMethod = clazz.methodsByName("main").get(0);
89         BreakpointRequest mainBp =
90             eventManager.createBreakpointRequest(mainMethod.location());
91         mainBp.enable();
92         System.out.println("Breakpoint set at main method start");
93     } catch (Exception e) {
94         System.err.println("Failed to set main breakpoint: " + e.getMessage());
95     }
96
97     // Встановлюємо breakpoint в методі fibonacci
98     try {
99         Method fibMethod = clazz.methodsByName("fibonacci").get(0);
100         BreakpointRequest fibBp =
101             eventManager.createBreakpointRequest(fibMethod.location());
102         fibBp.enable();
103         System.out.println("Breakpoint set in fibonacci method");
104     } catch (Exception e) {
105         System.err.println("Failed to set fibonacci breakpoint: " +
106             e.getMessage());
107     }
108 }
109
110 private void handleBreakpoint(BreakpointEvent event) {
111     try {
112         ThreadReference thread = event.thread();
113         StackFrame frame = thread.frame(0);
114         Location location = frame.location();
115         String methodName = location.method().name();

```

```

114     System.out.println("BREAKPOINT in " + methodName + "() at line " +
115         location.lineNumber());
116
117     if ("main".equals(methodName)) {
118         System.out.println("=== PROGRAM STARTED ===");
119         return;
120     }
121
122     if ("fibonacci".equals(methodName)) {
123         // Безпечноотримуюмопараметр n
124         Value nValue = null;
125         try {
126             List<LocalVariable> variables = location.method().variables();
127             for (LocalVariable var : variables) {
128                 if ("n".equals(var.name())) {
129                     nValue = frame.getValue(var);
130                     break;
131                 }
132             }
133         } catch (Exception varError) {
134             System.out.println("Cannot get variable 'n': " +
135                 varError.getMessage());
136         }
137
138         if (nValue != null) {
139             System.out.println("=== fibonacci(" + nValue + ") ===");
140         } else {
141             System.out.println("=== fibonacci(?) ===");
142         }
143
144         showStackFrames(thread);
145     }
146
147     } catch (Exception e) {
148         System.err.println("Error in breakpoint handler: " +
149             e.getClass().getSimpleName() + " - " + e.getMessage());
150         e.printStackTrace();
151     }
152 }
153
154 private void showStackFrames(ThreadReference thread) {
155     try {
156         List<StackFrame> frames = thread.frames();
157         int fibonacciCount = 0;
158
159         System.out.println("Stack depth: " + frames.size());
160
161         for (int i = 0; i < Math.min(frames.size(), 10); i++) {
162             StackFrame frame = frames.get(i);
163             Location location = frame.location();
164             String methodName = location.method().name();
165
166             if ("fibonacci".equals(methodName)) {
167                 // Безпечноотримуюмозмінну n
168                 String nValue = "?";
169                 try {
170                     List<LocalVariable> variables = location.method().variables();
171                     for (LocalVariable var : variables) {
172                         if ("n".equals(var.name())) {

```



```

170         Value value = frame.getValue(var);
171         if (value != null) {
172             nValue = value.toString();
173         }
174         break;
175     }
176 }
177 } catch (Exception e) {
178     // Ігноруємо помилки отримання змінних
179 }
180
181     System.out.println(" [" + i + "] fibonacci(n=" + nValue + ")");
182     fibonacciCount++;
183 } else {
184     System.out.println(" [" + i + "] " + methodName + "()");
185 }
186 }
187
188 if (frames.size() > 10) {
189     System.out.println(" ... and " + (frames.size() - 10) + " more
190     frames");
191 }
192
193 System.out.println("Fibonacci calls in stack: " + fibonacciCount);
194 System.out.println();
195 } catch (Exception e) {
196     System.err.println("Error showing stack: " + e.getClass().getSimpleName()
197     + " - " + e.getMessage());
198 }
199 }

```

Listing 4: Minimal Debugger

And target :

```

1 public class FibonacciTarget {
2     public static void main(String[] args){
3         System.out.println("Starting Fibonacci calculation...");
4         int n = 5;
5         long result = fibonacci(n);
6
7         System.out.println("fibonacci(" + n + ") = " + result);
8     }
9
10    public static long fibonacci(int n) {
11        if (n <= 1) {
12            return n;
13        }
14        return fibonacci(n - 1) + fibonacci(n - 2);
15    }
16 }

```

Listing 5: Debugger Target

We should compile it with debugging info :

```

1 javac -g -cp .;%JAVA_HOME%/lib/tools.jar *.java

```

To run app in debug mode we use this:

```
1 java -cp .;%JAVA_HOME%/lib/tools.jar  
-agentlib:jdwp=transport=dt_socket,server=y,suspend=y,address=5005 FibonacciTarget
```

Then run debugger :

```
1 java -cp .;%JAVA_HOME%/lib/tools.jar MinimalDebugger
```

Exploring debugger output :

```
1 java -cp .;C:\server\jdk-22_windows-x64_bin\jdk-22.0.1\lib/tools.jar MinimalDebugger  
2  
3 Connecting to target JVM...  
4 Connected to: Java HotSpot(TM) 64-Bit Server VM  
5 Class loaded: FibonacciTarget  
6 Breakpoint set at main method start  
7 Breakpoint set in fibonacci method  
8 BREAKPOINT in main() at line 3  
9 === PROGRAM STARTED ===  
10 BREAKPOINT in fibonacci() at line 11  
11 === fibonacci(5) ===  
12 Stack depth: 2  
13 [0] fibonacci(n=5)  
14 [1] main()  
15 Fibonacci calls in stack: 1  
16  
17 BREAKPOINT in fibonacci() at line 11  
18 === fibonacci(4) ===  
19 Stack depth: 3  
20 [0] fibonacci(n=4)  
21 [1] fibonacci(n=5)  
22 [2] main()  
23 Fibonacci calls in stack: 2  
24  
25 BREAKPOINT in fibonacci() at line 11  
26 === fibonacci(3) ===  
27 Stack depth: 4  
28 [0] fibonacci(n=3)  
29 [1] fibonacci(n=4)  
30 [2] fibonacci(n=5)  
31 [3] main()  
32 Fibonacci calls in stack: 3  
33  
34 BREAKPOINT in fibonacci() at line 11  
35 === fibonacci(2) ===  
36 Stack depth: 5  
37 [0] fibonacci(n=2)  
38 [1] fibonacci(n=3)  
39 [2] fibonacci(n=4)  
40 [3] fibonacci(n=5)  
41 [4] main()  
42 Fibonacci calls in stack: 4  
43  
44 BREAKPOINT in fibonacci() at line 11  
45 === fibonacci(1) ===  
46 Stack depth: 6  
47 [0] fibonacci(n=1)  
48 [1] fibonacci(n=2)  
49 [2] fibonacci(n=3)  
50 [3] fibonacci(n=4)
```

```

51     [4] fibonacci(n=5)
52     [5] main()
53 Fibonacci calls in stack: 5
54
55 BREAKPOINT in fibonacci() at line 11
56 === fibonacci(0) ===
57 Stack depth: 6
58     [0] fibonacci(n=0)
59     [1] fibonacci(n=2)
60     [2] fibonacci(n=3)
61     [3] fibonacci(n=4)
62     [4] fibonacci(n=5)
63     [5] main()
64 Fibonacci calls in stack: 5
65
66 BREAKPOINT in fibonacci() at line 11
67 === fibonacci(1) ===
68 Stack depth: 5
69     [0] fibonacci(n=1)
70     [1] fibonacci(n=3)
71     [2] fibonacci(n=4)
72     [3] fibonacci(n=5)
73     [4] main()
74 Fibonacci calls in stack: 4
75
76 BREAKPOINT in fibonacci() at line 11
77 === fibonacci(2) ===
78 Stack depth: 4
79     [0] fibonacci(n=2)
80     [1] fibonacci(n=4)
81     [2] fibonacci(n=5)
82     [3] main()
83 Fibonacci calls in stack: 3
84
85 BREAKPOINT in fibonacci() at line 11
86 === fibonacci(1) ===
87 Stack depth: 5
88     [0] fibonacci(n=1)
89     [1] fibonacci(n=2)
90     [2] fibonacci(n=4)
91     [3] fibonacci(n=5)
92     [4] main()
93 Fibonacci calls in stack: 4
94
95 BREAKPOINT in fibonacci() at line 11
96 === fibonacci(0) ===
97 Stack depth: 5
98     [0] fibonacci(n=0)
99     [1] fibonacci(n=2)
100    [2] fibonacci(n=4)
101    [3] fibonacci(n=5)
102    [4] main()
103 Fibonacci calls in stack: 4
104
105 BREAKPOINT in fibonacci() at line 11
106 === fibonacci(3) ===
107 Stack depth: 3
108     [0] fibonacci(n=3)
109     [1] fibonacci(n=5)

```

```

110     [2] main()
111 Fibonacci calls in stack: 2
112
113 BREAKPOINT in fibonacci() at line 11
114 === fibonacci(2) ===
115 Stack depth: 4
116     [0] fibonacci(n=2)
117     [1] fibonacci(n=3)
118     [2] fibonacci(n=5)
119     [3] main()
120 Fibonacci calls in stack: 3
121
122 BREAKPOINT in fibonacci() at line 11
123 === fibonacci(1) ===
124 Stack depth: 5
125     [0] fibonacci(n=1)
126     [1] fibonacci(n=2)
127     [2] fibonacci(n=3)
128     [3] fibonacci(n=5)
129     [4] main()
130 Fibonacci calls in stack: 4
131
132 BREAKPOINT in fibonacci() at line 11
133 === fibonacci(0) ===
134 Stack depth: 5
135     [0] fibonacci(n=0)
136     [1] fibonacci(n=2)
137     [2] fibonacci(n=3)
138     [3] fibonacci(n=5)
139     [4] main()
140 Fibonacci calls in stack: 4
141
142 BREAKPOINT in fibonacci() at line 11
143 === fibonacci(1) ===
144 Stack depth: 4
145     [0] fibonacci(n=1)
146     [1] fibonacci(n=3)
147     [2] fibonacci(n=5)
148     [3] main()
149 Fibonacci calls in stack: 3
150
151 Target VM terminated

```

Listing 6: Call frames in Java Debugger

Debugger class diagram shown here

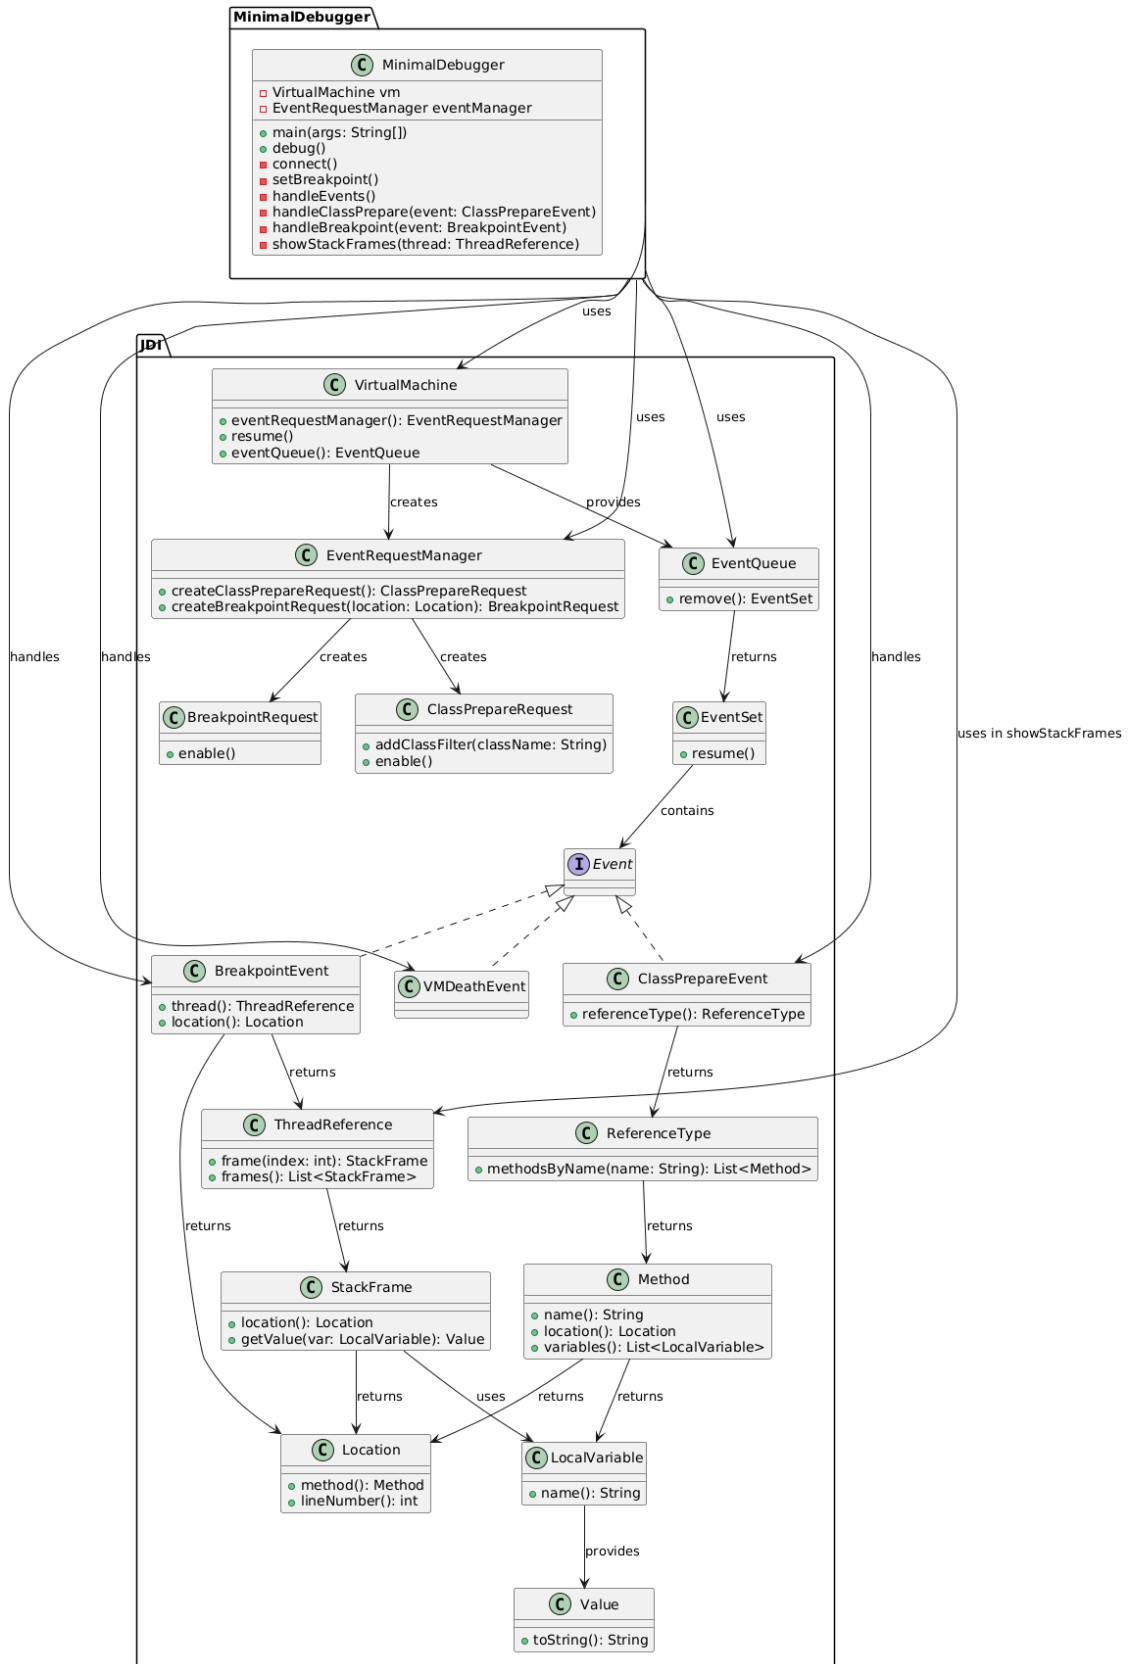


Figure 5: Debugger class diagram

2.11 Conclusion

The recursive Fibonacci algorithm has:

- **Time Complexity:** $O(\phi^n)$, where $\phi \approx 1.618$, due to the exponential number of nodes in the recursion tree.
- **Space Complexity:** $O(n)$, due to the linear depth of the call stack, as only one path of the recursion tree is active at any time.

3 Optimizing Recursion

3.1 Memoization

Memoization stores computed values to avoid redundant calculations.

Time Complexity: $T(n) = \mathcal{O}(n)$

Space Complexity: $M(n) = \mathcal{O}(n)$

```
1 import java.util.HashMap;
2 import java.util.Map;
3
4 private static final Map<Integer, Integer> memo = new HashMap<>();
5 public static int fibMemo(int n) {
6     if (n <= 1) return n;
7     if (memo.containsKey(n)) return memo.get(n);
8     int result = fibMemo(n - 1) + fibMemo(n - 2);
9     memo.put(n, result);
10    return result;
11 }
```

Listing 7: Memoized Fibonacci

3.2 Dynamic Programming

Dynamic programming (DP) avoids recursion entirely. We present two variants: array-based and space-optimized.

```
1 public static int fibDPArray(int n) {
2     if (n <= 0) return 0;
3     int[] fib = new int[n + 1];
4     fib[0] = 0;
5     fib[1] = 1;
6     for (int i = 2; i <= n; i++) {
7         fib[i] = fib[i - 1] + fib[i - 2];
8     }
9     return fib[n];
10 }
```

Listing 8: Array-Based DP Fibonacci

```
1 public static int fibDPOptimized(int n) {
2     if (n <= 0) return 0;
3     if (n == 1) return 1;
4     int prev = 0, curr = 1;
5     for (int i = 2; i <= n; i++) {
6         int next = prev + curr;
7         prev = curr;
8         curr = next;
9     }
10    return curr;
11 }
```

Listing 9: Space-Optimized DP Fibonacci

Time Complexity: $T(n) = \mathcal{O}(n)$

Space Complexity: $M(n) = \mathcal{O}(1)$ (optimized version)

3.3 Linear algebra and Fibonacci

3.3.1 Matrices and Transformations

We begin with the general concepts of matrices as linear transformations and matrix multiplication, then apply these ideas to the Fibonacci sequence, deriving its matrix form and geometric meaning.

Matrices represent linear transformations in vector spaces. A linear transformation $T : \mathbb{R}^m \rightarrow \mathbb{R}^n$ preserves vector addition and scalar multiplication:

$$T(\mathbf{u} + \mathbf{v}) = T(\mathbf{u}) + T(\mathbf{v}), \quad T(c\mathbf{u}) = cT(\mathbf{u}).$$

For example, a 2x2 matrix defines a transformation $T : \mathbb{R}^2 \rightarrow \mathbb{R}^2$. Geometrically, such transformations can include rotations, scalings, shears, or reflections in the plane.

In the context of sequences like Fibonacci, we will use a specific matrix to model the recurrence as a linear transformation that iteratively evolves a state vector.

Matrix multiplication is defined such that the product AB represents the composition of the linear transformations corresponding to B followed by A . For two matrices A (an $m \times n$ matrix) and B (an $n \times p$ matrix), the element $(AB)_{ij}$ is the dot product of the i -th row of A and the j -th column of B :

$$(AB)_{ij} = \sum_{k=1}^n a_{ik}b_{kj}.$$

This definition arises from the need to compose linear transformations. If T_A is the transformation for A and T_B for B , then $T_A(T_B(\mathbf{v})) = (AB)\mathbf{v}$. Improve this def !!!

Geometrically, matrix multiplication corresponds to applying one transformation after another. For example, multiplying rotation matrices composes rotations; scaling followed by shearing distorts shapes accordingly. In sequences, repeated multiplication evolves the state over multiple steps, leading to growth or convergence patterns observable in geometric structures.

Building on linear transformations, we express the Fibonacci recurrence in matrix form. Represent a pair of consecutive Fibonacci numbers as a vector:

$$\begin{bmatrix} F_n \\ F_{n-1} \end{bmatrix}.$$

We seek a matrix $A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$ such that:

$$A \cdot \begin{bmatrix} F_n \\ F_{n-1} \end{bmatrix} = \begin{bmatrix} F_{n+1} \\ F_n \end{bmatrix}.$$

Multiplying matrix to vector gives us the system of equations:

$$\begin{cases} aF_n + bF_{n-1} = F_{n+1} \\ cF_n + dF_{n-1} = F_n \end{cases} \quad (1)$$

From the Fibonacci recurrence, we know:

$$F_{n+1} = F_n + F_{n-1}.$$

Comparing this with the first equation from system (1), we get:

$$aF_n + bF_{n-1} = F_n + F_{n-1}.$$

For this to hold for all n , the coefficients must match:

$$a = 1, \quad b = 1.$$

For second equation from (1), we need:

$$cF_n + dF_{n-1} = F_n.$$

This implies:

$$c = 1, \quad d = 0.$$

Thus, the matrix is:

$$A = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}.$$

This matrix acts as a linear transformation:

$$A \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x + y \\ x \end{bmatrix},$$

which is a shear transformation.

To find later terms, we use matrix powers:

$$\begin{bmatrix} F_{n+1} \\ F_n \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}^n \cdot \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad (2)$$

This is what we were looking for - recurrence relation in matrix form for Fibonacci numbers. We'll visualize (2) later because it's fun.

3.3.2 Algorithm implementation in Java

```

1 public static long fibonacciMatrix(int n) {
2     if (n <= 0) return 0;
3     long[][] matrix = {{1, 1}, {1, 0}};
4     long[][] result = matrixPower(matrix, n - 1);
5     return result[0][0];
6 }
7
8 static long[][] matrixPower(long[][] matrix, int n) {
9     int row = matrix.length;
10    long[][] result = {{1, 0}, {0, 1}};
11    while (n > 0) {
12        if (n % 2 == 1) result = matrixMultiply(result, matrix);
13        matrix = matrixMultiply(matrix, matrix);
14        n /= 2;
15    }
16    return result;
17 }
18
19 static long[][] matrixMultiply(long[][] a, long[][] b) {
20     long[][] result = new long[2][2];
21     for (int i = 0; i < 2; i++)
22         for (int j = 0; j < 2; j++)
23             for (int k = 0; k < 2; k++)
24                 result[i][j] += a[i][k] * b[k][j];
25     return result;
26 }

```

Listing 10: Matrix Exponentiation for Fibonacci

3.3.3 Time and space complexity

Time Complexity: $T(n) = \mathcal{O}(\log n)$

Space Complexity: $M(n) = \mathcal{O}(1)$

3.4 Computer Algebra Systems

Using Maple, we can compute Fibonacci numbers efficiently:

```
1 with(combinat, fibonacci);
2 fibonacci_numbers := seq(fibonacci(i), i = 0 .. 100000):
3 f := fopen("fibonacci_numbers.txt", WRITE):
4 for num in fibonacci_numbers do
5   fprintf(f, "%a\n", num);
6 end do;
7 fclose(f);
```

Listing 11: Fibonacci in Maple

Run with:

```
1 cmaple -q fibonacci.mpl >out.log
```

Listing 12: Running Maple Script

3.5 Binet formula in real life

The Binet formula provides a closed-form expression for the n -th Fibonacci number:

$$F(n) = \frac{\phi^n - \psi^n}{\sqrt{5}},$$

where $\phi = \frac{1+\sqrt{5}}{2}$ (the golden ratio) and $\psi = \frac{1-\sqrt{5}}{2}$. Since it directly computes $F(n)$ without iterative steps, it has a theoretical time complexity of $O(1)$ for a single computation, assuming arithmetic operations (like exponentiation) are constant-time. However, despite this apparent efficiency, the Binet formula is impractical for large n in computational practice for several reasons.

1. **Numerical Instability with Floating-Point Arithmetic:** For large n , ϕ^n grows exponentially, while ψ^n (with $|\psi| < 1$) becomes extremely small. Their difference, $\phi^n - \psi^n$, involves subtracting two nearly equal values, leading to significant round-off errors in floating-point arithmetic (e.g., **double** in Java or C++). This causes inaccuracies, as seen when computing $F(71)$, where floating-point results may deviate (e.g., 308061521170131 instead of the correct 308061521170129).
2. **High Precision Requirements with Arbitrary-Precision Arithmetic:** To mitigate floating-point issues, arbitrary-precision arithmetic (e.g., **BigDecimal** in Java) can be used. However, computing ϕ^n , ψ^n , and $\sqrt{5}$ requires high precision (e.g., hundreds of decimal places for $n \approx 100$) to ensure the result is an exact integer after division by $\sqrt{5}$. This significantly increases computational cost, as each operation (exponentiation, division) scales with the number of digits, making the effective time complexity much worse than $O(1)$.
3. **Computational Overhead of Irrational Numbers:** The Binet formula involves irrational numbers (ϕ , ψ , $\sqrt{5}$), which are computationally expensive to handle with high precision. In contrast, iterative methods using integer arithmetic (e.g., with **BigInteger**) involve only simple additions, which are faster and inherently exact for Fibonacci numbers, as they are integers by definition.
4. **Scalability Issues for Large n :** For very large n (e.g., $n = 1000$), the precision required for $\sqrt{5}$ and ϕ^n grows proportionally to n , as the number of digits in $F(n)$ is approximately $n \cdot \log_{10}(\phi)$. This makes Binet's formula slower than iterative or matrix-based methods, which have logarithmic complexity ($O(\log n)$ with fast exponentiation) but operate on integers.
5. **Implementation Complexity:** Implementing Binet's formula requires careful management of precision and rounding to ensure integer results, which adds complexity compared to the straightforward iterative approach (e.g., $F(n) = F(n-1) + F(n-2)$). The latter is simpler to code and debug, as it avoids dealing with floating-point or arbitrary-precision libraries.

In conclusion, while the Binet formula is theoretically $O(1)$, its practical performance is hindered by numerical instability, high precision requirements, and computational overhead for irrational numbers. Iterative methods, or matrix exponentiation methods with $O(\log n)$ complexity, are preferred in practice due to their simplicity, exactness, and efficiency, especially for large n . For example, computing $F(71)$ iteratively with **BigInteger** is faster and guarantees the correct result (308061521170129) without precision management.

3.6 Conclusion

We explored Fibonacci computation methods, from naive recursion ($O(2^n)$) to matrix exponentiation ($O(\log n)$). Below is a comparison:

Method	Time Complexity	Space Complexity
Naive Recursion	$O(2^n)$	$O(n)$
Memoization	$O(n)$	$O(n)$
Dynamic Programming	$O(n)$	$O(1)$
Binet's Formula	$O(1)$	$O(1)$
Matrix Exponentiation	$O(\log n)$	$O(1)$

Table 4: Complexity Comparison of Fibonacci Methods

For large n , matrix exponentiation or BigDecimal-based DP are recommended. Future exploration could include fast doubling algorithms or parallel computation.

4 Limitations of Primitive Types for Large Fibonacci Numbers

4.1 Showcase : from int to double

```
1 import java.math.BigDecimal;
2 import java.util.HashMap;
3 import java.util.List;
4 import java.util.Map;
5
6 // Strategy interface for exact addition
7 interface AdditionStrategy<T extends Number> {
8     T add(T a, T b) throws ArithmeticException;
9 }
10
11 // Strategy for Integer
12 class IntegerAdditionStrategy implements AdditionStrategy<Integer> {
13     @Override
14     public Integer add(Integer a, Integer b) throws ArithmeticException {
15         return Math.addExact(a, b);
16     }
17 }
18
19 // Strategy for Long
20 class LongAdditionStrategy implements AdditionStrategy<Long> {
21     @Override
22     public Long add(Long a, Long b) throws ArithmeticException {
23         return Math.addExact(a, b);
24     }
25 }
26
27
28 // Strategy for Double
29 // Note: Unlike int/long, Java does not provide Math.addExact for double/float.
30 // This is because floating-point arithmetic (IEEE 754) does not throw overflow
31 // exceptions: instead, results silently become Infinity, -Infinity, or NaN.
32 // If strict overflow detection is required, it must be implemented manually
33 // (e.g., by checking Double.isInfinite / Double.isNaN), or by using BigDecimal
34 // for arbitrary precision arithmetic.
35
36 class FloatAdditionStrategy implements AdditionStrategy<Float> {
37     @Override
38     public Float add(Float a, Float b) throws ArithmeticException {
39         float result = a + b;
40
41         // Check for positive or negative infinity (overflow)
42         if (Float.isInfinite(result)) {
43             throw new ArithmeticException("Float overflow: " + a + " + " + b);
44         }
45
46         // Check for Not-a-Number result
47         if (Float.isNaN(result)) {
48             throw new ArithmeticException("Result is NaN: " + a + " + " + b);
49         }
50
51         return result;
52     }
53 }
54
```

```

55 class DoubleAdditionStrategy implements AdditionStrategy<Double> {
56     @Override
57     public Double add(Double a, Double b) throws ArithmeticException {
58         double result = a + b;
59         if (Double.isInfinite(result)) {
60             throw new ArithmeticException("Double overflow: " + a + " + " + b);
61         }
62         if (Double.isNaN(result)) {
63             throw new ArithmeticException("Result is NaN: " + a + " + " + b);
64         }
65         return result;
66     }
67 }
68
69 class FibonacciResult<T extends Number> {
70     private final int index;
71     private final T value;
72
73     private FibonacciResult(int index, T value) {
74         this.index = index;
75         this.value = value;
76     }
77
78
79     public static <T extends Number> FibonacciResult<T> of(int index, T value) {
80         return new FibonacciResult<>(index, value);
81     }
82
83     public Class<T> getType() {
84         return (Class<T>) value.getClass();
85     }
86
87     public int getIndex() {
88         return index;
89     }
90
91     public T getValue() {
92         return value;
93     }
94
95     public String toString() {
96         return switch (value) {
97             case Integer i -> Integer.toString(i);
98             case Long l -> Long.toString(l);
99             case Float f -> BigDecimal.valueOf(f).toPlainString();
100             case Double d -> BigDecimal.valueOf(d).toPlainString();
101
102             default -> value.toString();
103         };
104     }
105 }
106
107 @SuppressWarnings("all")
108 public class LargestExactFibonacci {
109     public static void main(String[] args) {
110         Map<Class<? extends Number>, AdditionStrategy<?>> strategies = new
111             HashMap<>();
112         strategies.put(Integer.class, new IntegerAdditionStrategy());
113         strategies.put(Long.class, new LongAdditionStrategy());

```

```

113 strategies.put(Float.class, new FloatAdditionStrategy());
114 strategies.put(Double.class, new DoubleAdditionStrategy());
115
116
117 FibonacciResult<Integer> resultInt = findLargestExactFibonacci(0, 1,
118     (AdditionStrategy<Integer>) strategies.get(Integer.class));
119 FibonacciResult<Long> resultLong = findLargestExactFibonacci(0L, 1L,
120     (AdditionStrategy<Long>) strategies.get(Long.class));
121 FibonacciResult<Float> resultFloat = findLargestExactFibonacci(0.0F, 1.0F,
122     (AdditionStrategy<Float>) strategies.get(Float.class));
123 FibonacciResult<Double> resultDouble = findLargestExactFibonacci(0.0D, 1.0D,
124     (AdditionStrategy<Double>) strategies.get(Double.class));
125
126 for (FibonacciResult<?> entry : List.of(resultInt, resultLong, resultFloat,
127     resultDouble)) {
128     System.out.printf("%s Fibonacci (%s): F(%s) = %s\n",
129         entry.getType().getSimpleName(), entry.getIndex(),
130         entry.getIndex(), entry);
131 }
132 private static <T extends Number> FibonacciResult<T> findLargestExactFibonacci(
133     T prev, T current, AdditionStrategy<T> strategy) {
134     int index = 1;
135     while (true) {
136         try {
137             T next = strategy.add(prev, current);
138             prev = current;
139             current = next;
140             index++;
141         } catch (ArithmeticException e) {
142             return FibonacciResult.of(index, current);
143         }
144     }
145 }
146 }

```

4.2 When double goes wild

4.2.1 $0.1 + 0.2 \neq 0.3$

In Java, when running the program

```

1 public class A {
2     public static void main(String[] args) {
3         double x = 0.1;
4         double y = 0.2;
5         double z = x + y;
6         System.out.println(z);
7     }
8 }

```

the output is

```
0.30000000000000004
```

instead of exactly 0.3.

If you think Java is broken, try it in JavaScript or Python
JS:

```
1 const x = 0.1;
2 const y = 0.2;
3 const z = x + y;
4 console.log(z);
```

Python:

```
1 x = 0.1
2 y = 0.2
3 z = x + y
4 print(z)
```

Output :

```
1 C:\temp\>java A
2 0.30000000000000004
3
4 C:\temp\>node a.js
5 0.30000000000000004
6
7 C:\temp\>python a.py
8 0.30000000000000004
```

This behavior is not a bug, but a direct consequence of the *IEEE 754 double-precision floating-point format*.

4.2.2 IEEE 754 Representation

A Java **double** uses 64 bits:

- 1 bit for the *sign*,
- 11 bits for the *exponent*,
- 52 bits for the *mantissa* (fraction).

The value is computed as

$$(-1)^{\text{sign}} \times 1.\text{mantissa} \times 2^{(\text{exponent}-1023)}.$$

4.2.3 Calculate as machines

The decimal fraction $0.1 = \frac{1}{10}$ has no finite binary expansion. In base 2 it becomes

$$0.1_{10} = 0.00011001100110011\dots_2$$

with repeating pattern **0011**. Since only 52 bits are available for the mantissa, this value is stored approximately as

$$0.1000000000000000000000055\dots$$

Similarly,

$$0.2 \approx 0.2000000000000000000000111, \quad 0.3 \approx 0.2999999999999999989.$$

Therefore,

$$0.1 + 0.2 \approx 0.3000000000000000000000444,$$

which is printed as **0.30000000000000004**.

The decimal system can exactly represent fractions like $1/10$ or $1/5$, but binary cannot. Binary floating-point can exactly represent only fractions whose denominators are powers of two (e.g., $1/2$, $1/4$, $1/8$). Numbers like $1/10$ or $1/5$ become infinite binary fractions, which must be rounded.

4.2.4 Conclusion

The result `0.30000000000000004` is not an error, but the inevitable effect of representing decimal fractions in binary with limited precision. Understanding IEEE 754 is essential for writing correct numerical programs.

4.3 Handling Large Numbers with `BigDecimal`

The `int` and `long` types in Java overflow for large Fibonacci numbers (e.g., F_{93} exceeds `long`). We use `BigDecimal` for arbitrary-precision arithmetic to address this limitation.

```
1 import java.math.BigDecimal;
2
3 public static BigDecimal fibBigDecimal(int n) {
4     if (n <= 0) return BigDecimal.ZERO;
5     if (n == 1) return BigDecimal.ONE;
6
7     BigDecimal prev = BigDecimal.ZERO;
8     BigDecimal curr = BigDecimal.ONE;
9     for (int i = 2; i <= n; i++) {
10         BigDecimal next = prev.add(curr);
11         prev = curr;
12         curr = next;
13     }
14     return curr;
15 }
```

Listing 13: Fibonacci with `BigDecimal`

Time Complexity: $T(n) = \mathcal{O}(n \cdot M)$, where M is the cost of `BigDecimal` operations.
Space Complexity: $M(n) = \mathcal{O}(1)$

5 Real-World Applications

5.1 Fibonacci heaps for Dijkstra's algorithm optimization

bla-bla-bla

5.2 Fibonacci retracement levels in stock market analysis

bla-bla-bla

5.3 Some of pseudorandom number generators

bla-bla-bla

References

- [1] Knuth, D. E. (1997). *The Art of Computer Programming, Volume 1: Fundamental Algorithms*. Addison-Wesley.
- [2] MapleSoft. (2025). Fibonacci Function Documentation. <https://www.maplesoft.com/support/help/Maple/view.aspx?path=combinat/fibonacci>.
- [3] Wikipedia. (2025). Fibonacci Number. https://en.wikipedia.org/wiki/Fibonacci_number.
- [4] SciPy docs https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.curve_fit.html